v0.5 — Reuse cannot save tokens without an approved shorter template

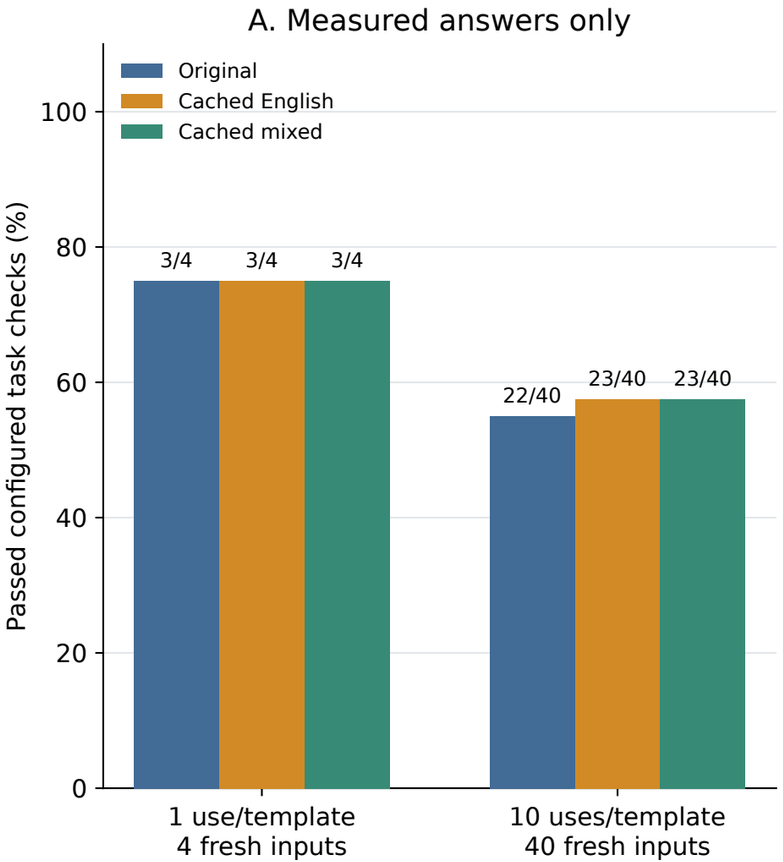

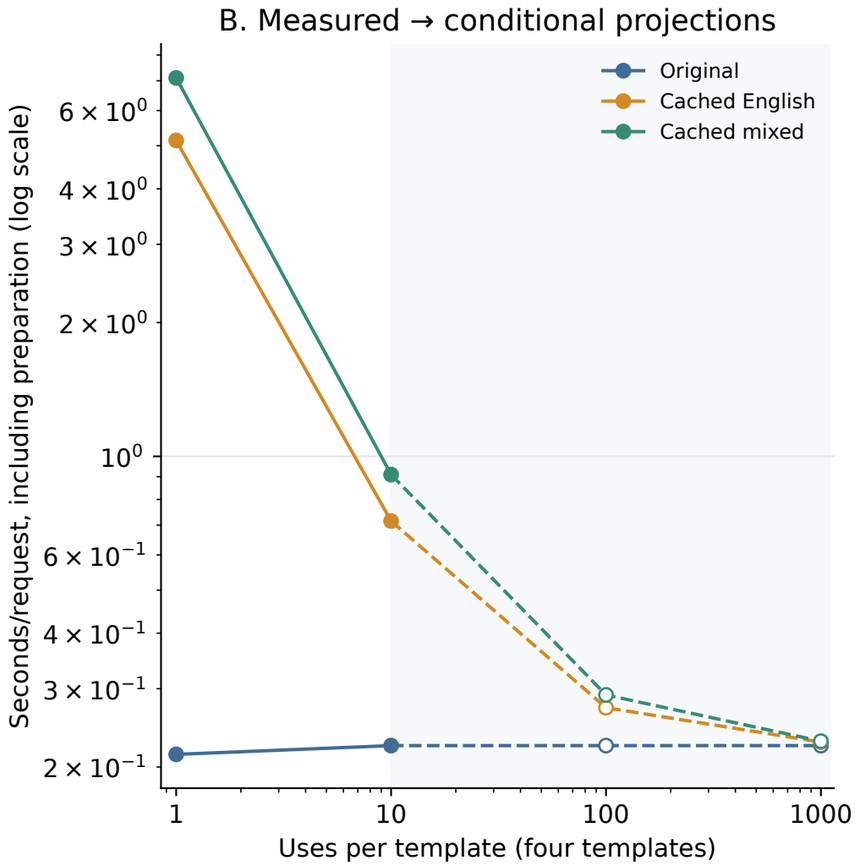

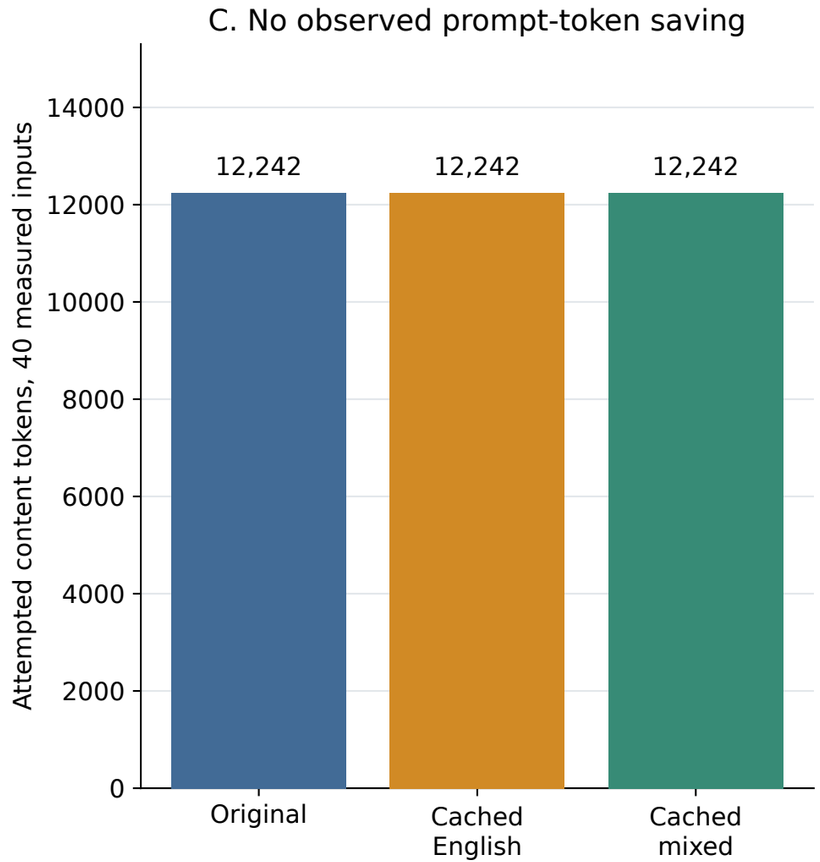

Four synthetic development templates, two calibration inputs each, then ten fresh inputs each. All test calls used original instructions.
Filled points: measured 1/10-use endpoints. Dashed/open points: 100/1000-use projections from observed mean times, not executed workloads.
Identical-prompt timing/answer variation is not a cache benefit. No accuracy forecast or monetary saving is claimed; tokens use o200k_base.
The 1/10-use accuracy endpoints have different case mixes; their difference does not establish an effect of reuse count.